



How to Win a Hackathon

From 40+ problems to a winning pitch — the process that judges can't ignore.

An agent that walks you through every step,
with 12 domain-expert judges reviewing your work.





How Most Teams Fail

They open a CSV of 40+ problems, pick one on gut feel, brainstorm 3 obvious ideas, and code for 47 hours.

What judges see: The wrong problem (nobody asked the owner). Zero validation (just the team's opinion). No awareness ("nobody else is doing this" — they are). A generic demo that works but tells no story.

Stage 1 — Problem Selection

Goal: Pick the right problem — not the first one or the easiest one.

Do 	Don't 
Cluster all problems by theme	Pick the one with the catchiest name
Score each cluster on impact × feasibility	Ignore the ones that sound boring
Check the market: competitors? already deployed?	Assume nobody has ever thought of this
Flag quality: too vague? needs hardware? too big?	Pretend a 6-month problem fits in 48 hours

The agent parses the CSV, clusters, scores, checks the market, and surfaces the top 3. You pick one.

Stage 2 — Problem Elicitation

Goal: Stop solving the wrong problem. Ask the owner.

Do 	Don't 
Ask: what hurts most <i>right now</i> ?	Assume the written statement is complete
Ask: what did you try that didn't work?	Rebuild something that already failed
Ask: who decides what gets bought?	Build a solution nobody can procure
Ask: what's the budget? timeline? environment?	Ignore real-world constraints
Let judges add hard questions	Trust only your perspective

The agent generates 6-10 questions. 5 judges add 2-3 domain-specific questions. In *real mode* it asks **you**. In *sim mode* it role-plays a persona. Re-score with the answers and **pick one**.

Stage 3 — Validation & Narrowing

Goal: Pick the right slice of the problem. Prove it's the right choice.

Do 	Don't 
Decompose into 5-8 sub-problems	Solve "the whole thing" in 48 hours
ROI-score each: impact, time-fit, demo-ability, risk	Pick the hardest sub-problem to look impressive
Validate with 5 domain experts	Rely on your team's gut feel
Pick the highest-ROI slice for 48 hours	Try to do everything

The agent decomposes into tractable sub-problems, scores each on 4 ROI axes, and every judge scores independently. Convergent picks = strong signal. You pick the slice.

Stage 4 — Solution Ideation

Goal: Generate 20+ ideas before picking one. The obvious solution is rarely the best.

Technique	Example
SCAMPER	Combine two existing solutions in a novel way
What would X do?	How would Palantir solve this? Anduril? A 16-year-old?
10X version	Unlimited budget, data, time — what's the moonshot?
Anti-solution	The worst possible idea (reveals what to avoid)
Analogy transfer	How would a video game do this? A cooking recipe?

The agent uses all 7 techniques to generate 20+ raw ideas. Deduped. Panel rates each 1-5. Top 5 = your shortlist.

Stage 5 — Solution Ranking

Goal: Pick the solution with the best shot at winning — not the one your team likes most.

Do 	Don't 
Web research every candidate: prior art, SOTA, competitors, TRL	Assume nobody has tried this
Re-rank with judges — weighted Borda, 5 perspectives	Use only your team's opinion
Track spread — genuine disagreement is healthy	Panic if judges disagree
Owner validates the final pick	Skip the owner at the end

The agent does web research, panel re-ranks with biases applied, and the owner (real or sim) validates the pick. Dissents recorded — they strengthen your story.

The 12-Judge Panel

5 auto-selected domain experts review every stage. Explicit biases, pet peeves, hard questions.
No yes-men.

Judge	"What I'd ask..."
Viper — F-16 pilot, 2200 hrs	"What kills this at 3am?"
Tran — EW specialist, MITRE	"Show me the false alarm rate."
Whitfield — Defense PM	"Where's the procurement path?"
Mehta — CTO, ex-Palantir	"What's already deployed?"
Kovalenko — Ukrainian drone op	"Tested in the dirt?"
Volkov — Red-team adversary	"I'd defeat this in 6 months."

Plus procurement, ethics, VC, scaling, intel, UX.

`/edth-agent panel viper` — Chat with any judge.

What Comes After

Once you've ranked your solution, the agent produces the rest:

Deliverable	Contents
Demo plan	3-min script, 30s pitch, risk register
Q&A prep	8-12 judge questions with answers
Market	TAM/SAM/SOM, trends, buyer personas
Competition	3+ rivals, strengths, weaknesses, moat
Business model	Revenue, pricing, GTM, defensibility
8-slide deck	HTML + PDF + PPTX — arrow-key navigable

Install & Run

```
git clone <repo> && cd edth-prob-solution-deck
python3 -m venv .venv && source .venv/bin/activate
pip install -e ".[dev]"
npm install -g @marp-team/marp-cli # optional, for PDF
```

```
/edth-agent dry-run      # 30s smoke test, full pipeline, zero interaction
/edth-agent run          # start a real interactive run
/edth-agent validate     # check all artefacts for issues
```

**Select it. Elicit it. Validate it. Ideate it.
Rank it.**

Don't get to the demo and realize you solved the wrong problem.